

Tao's Equational Proof Challenge Accepted

Lydia Kondylidou¹, Jasmin Blanchette¹, and Marijn J.H. Heule²

¹ Ludwig-Maximilians-Universität München, Munich, Germany

{l.kondylidou,jasmin.blanchette}@lmu.de

² Carnegie Mellon University, Pittsburgh, United States

marijn@cmu.edu

Abstract. In the context of the Equational Theories Project, Terence Tao posed the challenge of finding alternatives to a complicated 62-step proof found by the Vampire superposition prover. We introduce a proof minimization tool called Krympa. Using a combination of brute force and heuristics, and exploiting both Vampire and the Twee equational prover, the tool reduces the 62-step proof to 20 steps, each corresponding to a rewrite. In an empirical evaluation, it also performs well on 1431 equational problems originating from the same project, reducing in particular a 151-step proof to only 10 steps.

Keywords: Theorem provers · Equational logic · Proof minimization.

1 Introduction

The Equational Theories Project [6], launched in September 2024 by Fields medalist Terence Tao, aims at exploring the relations between different equational theories of magmas. A *magma* is a basic algebraic structure consisting of a set equipped with a single binary operation \diamond closed on that set. The project's first phase, concluded in April 2025, focused on equational laws for magmas that contain at most four applications of \diamond .

The project uses the Lean [21] proof assistant to express proofs and counter-examples but depends on automatic theorem provers and other external tools. The problems explored in the project's first phase all fall within first-order logic's unit equality fragment: They consist of a \forall -quantified equation as the sole axiom and a \forall -quantified equation as the proof goal, or conjecture.

For the problem 650 \implies 448, where 650 denotes the axiom $\forall x, y, z. x = x \diamond (y \diamond ((z \diamond x) \diamond y))$ and 448 denotes the conjecture $\forall x, y, z. x = x \diamond (y \diamond (z \diamond (x \diamond z)))$, the Vampire [18] superposition prover found a particularly complex proof, with 62 inference steps, excluding clauseification and Skolemization. Given that the proof is unintelligible, Tao challenged the community to find “an alternate proof, by whatever means you wish—human, semi-automated, or automated” [27].

One idea could be to run a specialized equational prover, Twee [23], instead of Vampire, but this results in a very long, 137-step proof. Another approach would be to use Lean's automation, such as the `aesop` [19], `canonical` [22], `duper` [8], and `grind` [1] tactics and the LeanHammer [28], to reconstruct and compress

consecutive superposition steps, in the style of Sledgehammer’s structured proof reconstruction [5, Sect. 6.3]. This would yield a shorter and more high-level proof, in which each step may combine multiple rewrites. Our approach is orthogonal. Our working hypothesis is that Vampire’s 62-step proof, which emerged as the byproduct of a saturation process, is likely suboptimal. By mixing and matching proofs generated by different automatic provers, as proposed by Sutcliffe et al. [26], we hope to achieve a shorter, simpler proof.

We introduce Krympa, a tool that minimizes equational proofs by decomposing them into independently provable components and reassembling them into more concise, intelligible proofs. Specifically, starting from a Vampire-generated proof, the tool transforms it into a direct proof (Sect. 3) and analyzes its inferences to break it down into intermediate results that serve as candidate lemmas. Each of these lemmas is then proved independently using Vampire and Twee (Sect. 4), the two leading systems in the unit equality division of CASC 2025 [25]. The resulting proofs are then combined into a single proof using heuristics that favor shorter derivations (Sect. 5).

Given the 62-step Vampire proof of $650 \implies 448$, our tool produces a 20-step proof, where 13 steps are generated by Twee (Sect. 6). In a larger empirical evaluation, we applied the tool to 1431 provable implications from the Equational Theories Project and obtained positive results (Sect. 7). In particular, the tool reduced a 151-step Vampire proof to 10 steps.

Our tool is implemented in Rust, OCaml, and Python. Its source code is available at <https://github.com/kondylidou/Krympa>. The files associated with Tao’s challenge and our empirical evaluation data are also available online [17].

2 Background

We briefly review the Vampire and Twee automatic provers and their associated proof formats.

2.1 Vampire and Superposition Proofs

Vampire is a saturation-based theorem prover for first-order logic with equality based on the superposition calculus [4]. It implements highly optimized search strategies and data structures, and integrates techniques such as literal selection, term orders, redundancy elimination, strategy scheduling, and portfolios.

Superposition works on implicitly \forall -quantified clauses. A preprocessor performs clausification and Skolemization. For example, the axiom $\forall x. f(x) = g(x)$ is transformed into $f(x) = g(x)$, where x is a free variable, and the conjecture $\forall x. f(x) = g(x)$ is negated and transformed into $f(\text{sk}) \neq g(\text{sk})$, where sk is a Skolem constant. The objective is to derive the contradictory clause \perp . For the unit equality fragment, the calculus’s two relevant inference rules are as follows:

$$\frac{t \neq u}{\perp} \text{equality resolution} \quad \frac{t = t' \quad s[u] \bowtie s'}{\mu(s[t'] \bowtie s')} \text{superposition}$$

The equality resolution rule has one premise, $t \neq u$, one conclusion, \perp , and one side condition: that t and u are unifiable. The superposition rule has two premises and one conclusion. The \bowtie symbol denotes either $=$ or \neq throughout the rule. The $=$ and \neq operators are commutative; for example, the premise $t = t'$ can match the equation $f(a) = b$ either as is or as $b = f(a)$. The premises are assumed to have disjoint sets of variables, which can be achieved by renaming. Also in the rule, $s[]$ is a term with a hole, the terms $s[u]$ and $s[t']$ are obtained by filling the hole in $s[]$ with u and t' , and μ is a most general unifier of t and u . For example, the most general unifier of the terms $h(a, y)$ and $h(x, b)$ is $\{x \mapsto a, y \mapsto b\}$; applying it on both terms yields $h(a, b)$. Finally, the rule has further side conditions, not shown here, that restrict the search space.

Example 1. A subtle case of the superposition rule arises when both premises are the same clause. Consider the following rule instance, where the variable in the second premise has been renamed to avoid a clash:

$$\frac{f(f(x)) = g(x) \quad f(f(x')) \neq g(x')}{f(g(x)) \neq g(f(x))} \text{superposition}$$

This instance is obtained by taking $t := f(f(x))$, $t' := g(x)$, $\bowtie := \neq$, $s[] := f([])$, $u := f(x')$, $s' := g(x')$, and $\mu = \{x' \mapsto f(x)\}$. Applying the unifier μ to both premises yields the equation $f(f(x)) = g(x)$ and the disequation $f(f(f(x))) \neq g(f(x))$. The inference replaces the subterm $f(f(x))$ in the disequation with $g(x)$ using the equation as a left-to-right rewrite rule, and derives the conclusion. ■

Example 2. Vampire implements *parallel superposition*, a variant of the superposition rule in which multiple subterms that match a term are replaced. The following inference illustrates this:

$$\frac{b = a \quad h(b, a, b) \neq h(a, b, a)}{h(a, a, a) \neq h(a, a, a)} \text{parallel superposition} \quad \blacksquare$$

Superposition proofs are represented in a linear format. They are refutational and show how to derive \perp from the input axioms and the negated conjecture.

Example 3. The following is a linear superposition proof from clauses 1–3:

1. $a = b$	axiom
2. $f(x) = x$	axiom
3. $h(f(b), a) \neq h(a, f(b))$	negated conjecture
4. $h(b, a) \neq h(a, b)$	by parallel superposition from 2 and 3
5. $h(a, a) \neq h(a, a)$	by parallel superposition from 1 and 4
6. \perp	by equality resolution from 5

2.2 Twee and Structured Equational Chain Proofs

Twee is an automatic prover specialized for equational reasoning. It is based on the unfailing completion procedure [3], an extension of Knuth–Bendix completion [16]. In the DISCOUNT and Waldmeister tradition [7], Twee’s proofs are structured as a sequence of lemmas, where each lemma and the conjecture are

110 proved by a chain of equalities. Twee introduces lemmas if they are needed more
 111 than once. Twee proofs are arguably more readable than Vampire proofs. As
 112 with superposition, quantifiers are eliminated by a preprocessor.

113 **Example 4.** The following is a Twee-style proof of goal 1 from axioms 1 and 2:

$$\begin{array}{ll}
 \text{Axiom 1: } a = b & \text{Goal 1: } h(f(b), a) = h(a, f(b)) \\
 \text{Axiom 2: } f(x) = x & \text{Proof:} \\
 \text{Lemma 3: } f(b) = a & h(f(b), a) \\
 \text{Proof:} & = \{ \text{ by lemma 3 } \} \\
 f(b) & h(a, a) \\
 = \{ \text{ by axiom 1 right-to-left } \} & = \{ \text{ by lemma 3 right-to-left } \} \\
 f(a) & h(a, f(b)) \\
 = \{ \text{ by axiom 2 } \} & \\
 a & \blacksquare
 \end{array}$$

115 3 Proof Redirection

116 Vampire generates proofs by refutation, whereas our mix-and-match approach
 117 requires direct proofs. To bridge this gap, we transform Vampire proofs into
 118 direct proofs. In the following sections, we will always use direct proofs.

119 To redirect a proof by refutation in equational logic, we first introduce \exists
 120 quantifiers for Skolem constants and \forall quantifiers for variables. For example,
 121 $h(x, sk) \neq x$ is transformed into $\exists z. \forall x. h(x, z) \neq x$. Then we apply the contra-
 122 positive to all inferences in which a premise and the conclusion are disequations
 123 to obtain positive equations. Thus, the inference

$$\frac{h(a, y) = b \quad h(x, sk) \neq x}{b \neq a} \text{superposition}$$

124 becomes

$$\frac{\forall y. h(a, y) = b \quad b = a}{\forall z. \exists x. h(x, z) = x}$$

125 Equality resolution inferences from a premise $t \neq t$ are omitted since their con-
 126 trapositives derive trivial equations.

127 **Example 5.** The following is a direct proof obtained from Example 3's proof
 128 by refutation.

$$\begin{array}{ll}
 1. a = b & \text{axiom} \\
 2. \forall x. f(x) = x & \text{axiom} \\
 3. h(b, a) = h(a, b) & \text{from 1 and } h(a, a) = h(a, a) \\
 4. h(f(b), a) = h(a, f(b)) & \text{from 2 and 3} \blacksquare
 \end{array}$$

131 4 Proof Generation for Intermediate Lemmas

132 Our approach starts by translating the main theorem into a TPTP [12] input
 133 problem and running Vampire to produce an initial proof. This proof is turned

134 into a direct proof, then decomposed into intermediate lemmas. For each lemma,
 135 we generate corresponding problems, with the objective of proving them using
 136 Vampire and Twee. Three problem variants are generated:

- 137 1. *Big-step problems* contain the axioms together with the lemma as the conjecture,
 138 and nothing else. This allows us to investigate whether a radically new proof, with different intermediate steps, can be found.
- 140 2. *Small-step problems* contain the axioms together with the lemma as the conjecture,
 141 and all lemmas derived prior to this lemma in the initial proof as additional axioms.
 142 This allows us to investigate whether a somewhat similar variant of the original derivation can be found.
- 144 3. *Abstracted problems* are variants of big-step problems that contain the axioms
 145 together with an abstracted version of the lemma as the conjecture.
 146 Specifically, selected subterms of the lemma—for example, expressions such
 147 as $x \diamond y$ that do not contain nested applications—are replaced by fresh variables.
 148 This allows us to investigate whether a more general version of the lemma is provable, ideally with a shorter, more abstract proof.

150 Each problem is submitted to the two provers. If a proof is found for a small-step
 151 problem, we expand it to recursively include the shortest proofs of the lemmas
 152 used as axioms for the axioms referenced in the proof. Ties are broken arbitrarily.
 153 Note that abstracted problems might be unprovable.

154 Next, we compare the proofs of the three problem variants corresponding to
 155 the same lemma. If the abstracted problem has the shortest proof, the lemma
 156 it proves is replaced in all small-step problems where it appears as an axiom
 157 with the generalized lemma from the abstracted problem. Each updated small-
 158 step problem is then re-proved, and if the result has fewer steps, we replace the
 159 small-step problem's proof with it.

160 The length of a Vampire-generated proof is the number of steps of its redirected
 161 proof, excluding preprocessing. For Twee, the length of a proof is the
 162 cumulative number of equalities in the equality chains. Thus, the Vampire proof
 163 in Example 5 has two steps, and the Twee proof in Example 4 has four steps.

164 5 Proof Construction for the Main Theorem

165 Based on the intermediate lemmas' proofs generated in the previous phase, our
 166 approach constructs a proof of the main theorem. The proof generally consists
 167 of three segments. The first segment starts with the axioms and ends with the
 168 derivation of a so-called *departure lemma*. The second segment derives a so-called
 169 *arrival lemma*. The third segment derives the conjecture. Different candidates
 170 are considered as the departure and arrival lemmas, yielding different proofs.
 171 The proof with the fewest steps is chosen.

172 Specifically, we first identify up to six intermediate lemmas that arise close to
 173 the end of the initial proof, including the conjecture, and consider them as can-
 174 didate arrival lemmas. For each of these, we consider its transitive dependencies

as candidate departure lemmas. Then, for each candidate departure lemma, we construct a problem with the axioms and the departure lemma's dependencies as the axioms and the departure lemma itself as the conjecture. We run both provers and, if at least one succeeds, we use the shorter result as the proof of the first segment, unless an even shorter proof was generated in the previous phase.

Next, for each pair of candidate departure and arrival lemmas, we generate a new problem with the original axioms, the departure lemma, and its dependencies as axioms and the arrival lemma as the conjecture. We run both provers and, if at least one succeeds, we use the shorter result as the proof of the second segment, unless an even shorter proof was generated earlier. Finally, we generate a new problem with the original axioms, the departure lemma, its dependencies, and the arrival lemma as axioms and the original conjecture as the conjecture. We run both provers and, if at least one succeeds, we use the shorter result as the proof of the third segment, unless an even shorter proof was generated earlier.

Without the separation into segments, proof minimization could be intractable due to combinatorial explosion. We chose to work with three segments as a trade-off between performance and flexibility.

Example 6. Before we review the three-segment proof construction approach in detail, let us look at an example. The following sketch represents an initial seven-step Vampire-generated redirected proof of a theorem $A \implies C$:

195	A	axiom
196	L_1	from A and A
197	L_2	from A and L_1
198	L_3	from L_1 and L_2
199	L_4	from L_2 and L_3
200	L_5	from L_3 and L_4
201	L_6	from A and L_5
202	C	from L_5 and L_6

Here, A denotes the axiom, and L_1, \dots, L_6 are the lemmas used to derive the conjecture C .

In the first phase, for each lemma L_1, \dots, L_6 , we construct big-step, small-step, and abstracted problems and try to prove them using Vampire and Twee, retaining the shortest proof for each lemma. Suppose the following: The shortest proof of L_1 has one step and is obtained from its big-step problem using Vampire; for L_2 and L_3 , the shortest proofs are obtained from their small-step problems using Twee; for L_4 , the shortest proof is obtained from its abstracted problem using Twee; and for L_5 and L_6 , the shortest proofs are obtained from their small-step problems using Vampire.

In the next phase, the last five lemmas, L_2, \dots, L_6 , and the conjecture C are considered as candidate arrival lemmas. We focus on L_6 . The proof below, found by Vampire for L_6 's small-step problem, is the shortest proof for L_6 :

216	A	axiom
217	L_1	from A and A
218	L_2	from A and L_1

219	L_3	from L_1 and L_2
220	L_4	from L_2 and L_3
221	L_5	from L_3 and L_4
222	L_6	from A and L_5

223 This proof happens to be identical to the first six steps of the initial proof, but
 224 in general it could be different.

225 Next, lemmas L_1 to L_5 are considered as candidate departure lemmas. We
 226 focus on L_3 . The proof of conjecture C is constructed by concatenating three
 227 segments. For the first segment, we create a new problem with A , L_1 , and L_2
 228 as axioms, since they are dependencies of the departure lemma L_3 in the above
 229 proof of L_6 , and L_3 as the conjecture. We run both provers on this problem
 230 and obtain a two-step Vampire proof of L_3 from A , L_1 , and a new lemma L'_2 .
 231 Since L_1 is treated as an axiom, we must include its proof to obtain a complete
 232 proof of L_3 . In the first phase, we found a one-step Vampire proof of L_1 from
 233 the axiom A , so we use it. In summary, the proofs of L_1 and L_3 form the first
 234 segment, which consists of one step for L_1 and two steps for L_3 .

235 For the second segment, we create a new problem with A , L_1 , L'_2 , and L_3 as
 236 axioms and the arrival lemma L_6 as the conjecture. We run both provers on this
 237 problem and obtain a two-step Twee proof of L_6 from L_1 and L_3 . Together with
 238 the first segment, this yields a five-step proof of L_6 . Since this proof is shorter
 239 than the six-step proof of L_6 presented above, it is used as the second segment.

240 For the third segment, we create a new problem with A , L_1 , L'_2 , the departure
 241 lemma L_3 , and the arrival lemma L_6 as axioms and C as the conjecture. We run
 242 both provers on this problem and obtain a two-step Twee proof of C from L'_2
 243 and L_3 . Since this proof does not use the arrival lemma L_6 , the second segment
 244 is excluded from the result. Concatenating the first and third segments yields a
 245 new five-step proof of C :

246	A	axiom
	L_1	from A
	L'_2	from A and L_1
	L_3	from L_1 and L'_2
	C	by a two-step equality chain using L'_2 and L_3

247 Finally, other combinations of candidate departure and arrival lemmas are
 248 also considered, and the shortest proof is retained. ■

249 5.1 Construction of the Dependency Graph

250 We identify lemmas occurring close to the end of the derivation as candidate
 251 arrival lemmas. Different candidates typically depend on substantially different
 252 subsets of earlier lemmas. Each candidate therefore induces its own dependency
 253 chain, and different choices can lead to substantially different proof lengths. We
 254 consider six candidate arrival lemmas extracted from the initial proof, including
 255 the conjecture itself, since our approach may produce a shorter proof of the
 256 conjecture by reproving it directly from a minimized dependency set.

257 For every candidate, we build a dependency graph that captures the lemmas
 258 required to derive it. Dependencies are determined from the shortest Vampire
 259 or Twee proof obtained for each lemma. Given that we generate three problem
 260 variants and run two provers, up to six proofs per lemma are considered. A
 261 lemma ℓ is considered to directly depend on a lemma ℓ' if the shortest proof
 262 of ℓ uses ℓ' as an axiom. Thus, for big-step and abstracted problems, only the
 263 original axioms can be dependencies. For small-step problems, each intermediate
 264 step in a Vampire proof and each lemma in a Twee proof is considered a lemma.

265 The dependency graph associated with a candidate arrival lemma is a di-
 266 rected acyclic graph (DAG) whose nodes correspond to lemmas and whose edges
 267 express derivability between them. Formally, let V be a finite set of lemmas, each
 268 represented by an equation and a set of dependencies on other lemmas. We con-
 269 struct a DAG (V, E) , where each vertex $\ell \in V$ corresponds to a lemma and each
 270 edge $(\ell, \ell') \in E$ indicates that lemma ℓ directly depends on lemma ℓ' . As an
 271 optimization, we merge lemmas that are identical up to the naming of variables,
 272 keeping the shortest proof.

273 5.2 Construction of the First Proof Segment

274 For each candidate arrival lemma, we investigate whether all lemmas included in
 275 its dependency graph are needed to derive it or whether a shorter proof can be
 276 obtained by choosing a departure lemma and recomputing parts of the derivation
 277 by combining proofs generated by the provers.

278 As candidate departure lemmas, we consider all lemmas in the DAG. Let ℓ
 279 be a candidate departure lemma. If ℓ depends only on the axioms, we take the
 280 shortest big-step, small-step, or abstracted proof previously found by Vampire or
 281 Twee. Otherwise, we build a problem that includes ℓ 's dependencies in the DAG
 282 as axioms and the departure lemma as the conjecture, and we run Vampire and
 283 Twee. If at least one of them succeeds, we choose the shorter proof as ℓ 's proof.
 284 This derivation, together with the shortest proofs of ℓ 's dependencies generated
 285 for the big-step, small-step, or abstracted problems, forms the first segment of
 286 the final proof. However, if we found an even shorter proof for the big-step, small-
 287 step, or abstracted problem, we use that proof instead. For small-step proofs, we
 288 must also include the proofs of the intermediate lemmas encoded as axioms.

289 5.3 Construction of the Remaining Proof Segments

290 To construct the second segment, we generate a problem with the departure
 291 lemma and its dependencies as axioms and the arrival lemma as the conjecture,
 292 and run both provers. If at least one of them succeeds, we choose the shorter
 293 proof as the proof of the arrival lemma. As above, we fall back on the proof of
 294 a big-step, small-step, or abstracted problem if it is even shorter.

295 Finally, to construct the third segment, we generate a problem with the
 296 departure lemma, the arrival lemma, and their dependencies as axioms and the
 297 original conjecture as the conjecture, and invoke both provers. If at least one of

them succeeds, we choose the shorter proof as the proof of the original conjecture.
As above, we fall back on a previously derived proof if it is even shorter.

The final proof is obtained by concatenating the three segments. The proof might contain unreferenced lemmas; these are pruned.

5.4 Proof Output

Our tool generates the minimized proof in a native format, from which two Lean outputs are produced. The first Lean output is a step-by-step formalization using the `calc` tactic to reconstruct chains of equalities. It applies the `duper` tactic to fill in the subproofs. For example, a proof of $t_1 = t_2 = t_3 = t_4$ would be represented by

```
calc
t1 = t2 := by duper ...
_ = t3 := by duper ...
_ = t4 := by duper ...
```

where the ellipses stand for `duper`'s arguments. The second Lean output is a more compact Lean formalization in which each lemma is proved directly using Lean's automation without including the intermediate steps in chains of equalities.

6 Application to Tao's Challenge

We implemented our approach and tried the resulting tool, Krympa, on Tao's challenge theorem 650 \implies 448:

$$(\forall x, y, z. x = x \diamond (y \diamond ((z \diamond x) \diamond y))) \implies \forall x, y, z. x = x \diamond (y \diamond (z \diamond (x \diamond z))).$$

Our tool first ran Vampire to obtain an initial 62-step superposition proof. Then it constructed 62 problems of each variant (big-step, small-step, and abstracted) and tried to prove them using Vampire and Twee. Among the six candidate arrival lemmas, the shortest proof was found by selecting

$$\forall x, y, z. x = x \diamond ((y \diamond ((z \diamond y) \diamond y)) \diamond x). \quad (\text{lemma 9})$$

The coloring highlights repeating patterns. Next, our tool constructed the dependency graph for this lemma. The DAG contained 37 lemmas. It was based on big- and small-step proofs.

Among the 37 candidate departure lemmas, our tool found the shortest proof by selecting

$$\begin{aligned} \forall x, y, z, w. & (x \diamond ((y \diamond x) \diamond x)) \diamond z = \\ & ((x \diamond ((y \diamond x) \diamond x)) \diamond z) \diamond (w \diamond ((x \diamond ((y \diamond x) \diamond x)) \diamond w)). \end{aligned} \quad (\text{lemma 7})$$

According to the DAG, the shortest proof of this lemma was found by running Vampire on the small-step problem consisting of the axiom and the following lemma dependencies:

$$\begin{aligned}
& \forall x, y, z, w. x \diamond ((y \diamond z) \diamond x) = \\
& \quad (x \diamond ((y \diamond z) \diamond x)) \diamond (w \diamond (z \diamond w)) && \text{(lemma 1)} \\
& \forall x, y, z, w, v, u. x \diamond ((y \diamond ((z \diamond w) \diamond y)) \diamond x) = \\
& \quad (x \diamond ((y \diamond ((z \diamond w) \diamond y)) \diamond x)) \diamond (v \diamond ((u \diamond (w \diamond u)) \diamond v)) && \text{(lemma 2)} \\
& \forall x, y, z, w, v. x \diamond (y \diamond x) = \\
& \quad (x \diamond (y \diamond x)) \diamond (z \diamond ((w \diamond ((v \diamond y) \diamond w)) \diamond z)) && \text{(lemma 3)} \\
& \forall x, y, z, w, v. x \diamond (y \diamond x) = \\
& \quad (x \diamond (y \diamond x)) \diamond ((z \diamond (y \diamond z)) \diamond (w \diamond ((v \diamond y) \diamond w))) && \text{(lemma 4)} \\
& \forall x, y, z, w. x \diamond ((y \diamond ((z \diamond y) \diamond y)) \diamond x) = \\
& \quad (x \diamond ((y \diamond ((z \diamond y) \diamond y)) \diamond x)) \diamond \\
& \quad (w \diamond ((y \diamond ((z \diamond y) \diamond y)) \diamond w)) && \text{(lemma 5)} \\
& \forall x, y, z, w. (x \diamond ((y \diamond x) \diamond x)) \diamond z = \\
& \quad ((x \diamond ((y \diamond x) \diamond x)) \diamond z) \diamond ((w \diamond ((x \diamond ((y \diamond x) \diamond x)) \diamond w)) \diamond \\
& \quad (z \diamond ((x \diamond ((y \diamond x) \diamond x)) \diamond z))). && \text{(lemma 6)}
\end{aligned}$$

Following the inference steps of the initial Vampire proof, our tool derived lemma 1 by applying a superposition inference with the axiom $x = x \diamond (y \diamond ((z \diamond x) \diamond y))$ as the first premise and a renamed copy $x' = x' \diamond (y' \diamond ((z' \diamond x') \diamond y'))$ as the second premise. The most general unifier of the first premise's right-hand side and the subterm $z' \diamond x'$ of the second premise is $\{x' \mapsto y \diamond ((z \diamond x) \diamond y), z' \mapsto x\}$. Applying the unifier to both premises yields the equations $x = x \diamond (y \diamond ((z \diamond x) \diamond y))$ and $y \diamond ((z \diamond x) \diamond y) = (y \diamond ((z \diamond x) \diamond y)) \diamond (y' \diamond ((x \diamond (y \diamond ((z \diamond x) \diamond y))) \diamond y'))$. The superposition inference replaced the subterm $x \diamond (y \diamond ((z \diamond x) \diamond y))$ in the second equation with x using the first equation as a right-to-left rewrite rule, and thus derived lemma 1, up to the naming of variables. Lemmas 2 to 7 were derived similarly following the steps of the initial Vampire proof.

Next, from the axiom and lemma 7, our tool proved the arrival lemma (lemma 9) using Twee. For this proof, Twee introduced the auxiliary lemma

$$\begin{aligned}
& \forall x, y, z, w. (y \diamond ((z \diamond y) \diamond y)) \diamond w = \\
& \quad ((y \diamond ((z \diamond y) \diamond y)) \diamond w) \diamond ((y \diamond ((z \diamond y) \diamond y)) \diamond x). && \text{(lemma 8)}
\end{aligned}$$

Finally, assuming all the lemmas derived so far, our tool proved the conjecture from lemmas 5 and 9 using Twee. The resulting proof has 20 steps, including three Twee-generated chains of equalities.

Below we present the final proof adapted from our tool's detailed Lean output. Instead of relying on proof automation, we use the `nth_rw` tactic, which performs a single rewrite step, where the numeric argument indicates which matching occurrence should be rewritten. In one case, two numbers are supplied, corresponding to a parallel rewrite.

```

351 class Magma (α : Type _) where
352   op : α → α → α
353
354   infix:65 " ◇ " => Magma.op

```

```

355 theorem Equation650_implies_Equation448 (G : Type _) [Magma G]
356   (op_law : ∀ x y z : G, x = x ◊ (y ◊ ((z ◊ x) ◊ y))) :
357   ∀ x y z : G, x = x ◊ (y ◊ (z ◊ (x ◊ z))) := 
358   have lemma1 (x y z w : G) :
359     x ◊ ((y ◊ z) ◊ x) = (x ◊ ((y ◊ z) ◊ x)) ◊ (w ◊ (z ◊ w)) := by
360   nth_rw 3 [op_law z x y]
361   exact op_law (x ◊ ((y ◊ z) ◊ x)) w z
362
363   have lemma2 (x y z w v u : G) :
364     x ◊ ((y ◊ ((z ◊ w) ◊ y)) ◊ x) =
365     (x ◊ ((y ◊ ((z ◊ w) ◊ y)) ◊ x)) ◊ (v ◊ ((u ◊ (w ◊ u)) ◊ v)) := by
366   nth_rw 1 2 [lemma1 y z w u]
367   exact lemma1 x (y ◊ ((z ◊ w) ◊ y)) (u ◊ (w ◊ u)) v
368
369   have lemma3 (x y z w v : G) :
370     x ◊ (y ◊ x) = (x ◊ (y ◊ x)) ◊ (z ◊ ((w ◊ ((v ◊ y) ◊ w)) ◊ z)) := by
371   nth_rw 1 [lemma1 w v y x]
372   exact op_law (x ◊ (y ◊ x)) z (w ◊ ((v ◊ y) ◊ w))
373
374   have lemma4 (x y z w v : G) :
375     x ◊ (y ◊ x) = (x ◊ (y ◊ x)) ◊ ((z ◊ (y ◊ z)) ◊ (w ◊ ((v ◊ y) ◊ w))) := by
376   nth_rw 1 [lemma1 w v y z]
377   exact lemma3 x y (z ◊ (y ◊ z)) w v
378
379   have lemma5 (x y z w : G) :
380     x ◊ ((y ◊ ((z ◊ y) ◊ y)) ◊ x) =
381     (x ◊ ((y ◊ ((z ◊ y) ◊ y)) ◊ x)) ◊ (w ◊ ((y ◊ ((z ◊ y) ◊ y)) ◊ w)) := by
382   nth_rw 1 [lemma2 w y z y x ((z ◊ y) ◊ y)]
383   exact lemma4 x (y ◊ ((z ◊ y) ◊ y)) w x ((z ◊ y) ◊ y)
384
385   have lemma6 (x y z w : G) :
386     (x ◊ ((y ◊ x) ◊ x)) ◊ z =
387     ((x ◊ ((y ◊ x) ◊ x)) ◊ z) ◊ ((w ◊ ((x ◊ ((y ◊ x) ◊ x)) ◊ w)) ◊
388     (z ◊ ((x ◊ ((y ◊ x) ◊ x)) ◊ z))) := by
389   nth_rw 1 [lemma5 z x y w]
390   exact op_law ((x ◊ ((y ◊ x) ◊ x)) ◊ z) (w ◊ ((x ◊ ((y ◊ x) ◊ x)) ◊ w)) z
391
392   have lemma7 (x y z w : G) :
393     (x ◊ ((y ◊ x) ◊ x)) ◊ z =
394     ((x ◊ ((y ◊ x) ◊ x)) ◊ z) ◊ (w ◊ ((x ◊ ((y ◊ x) ◊ x)) ◊ w)) := by
395   nth_rw 1 [lemma5 w x y z]
396   exact lemma6 x y z w
397
398   have lemma8 (x y z w : G) :
399     ((x ◊ ((y ◊ x) ◊ x)) ◊ z) ◊ ((x ◊ ((y ◊ x) ◊ x)) ◊ w) =
400     (x ◊ ((y ◊ x) ◊ x)) ◊ z := by
401   let T := x ◊ ((y ◊ x) ◊ x)
402   calc
403     (T ◊ z) ◊ (T ◊ w) =
404     ((T ◊ z) ◊ ((T ◊ w) ◊ ((T ◊ (T ◊ w)) ◊ ((w ◊ (T ◊ w)) ◊ (T ◊ (T ◊ w)))))) := by

```

```

405      nth_rw 1 [←op_law]
406      - = ((T ◊ z) ◊ ((T ◊ w) ◊ ((T ◊ (T ◊ w)) ◊ ((w ◊ (T ◊ w)) ◊
407          (T ◊ ((T ◊ w) ◊ (w ◊ (T ◊ w))))))) := by
408      nth_rw 1 [←lemma7]
409      - = ((T ◊ z) ◊ ((T ◊ w) ◊ ((T ◊ (T ◊ w)) ◊ ((w ◊ (T ◊ w)) ◊
410          (T ◊ ((T ◊ w) ◊ (w ◊ (T ◊ w)))) ◊ (((T ◊ w) ◊ (w ◊ (T ◊ w)) ◊
411              (T ◊ ((T ◊ w) ◊ (w ◊ (T ◊ w)))))))) := by
412      nth_rw 2 [←lemma7]
413      - = ((T ◊ z) ◊ ((T ◊ w) ◊ ((T ◊ (T ◊ w)) ◊ (w ◊ (T ◊ w))))) := by
414      nth_rw 1 [←op_law]
415      - = ((T ◊ z) ◊ ((T ◊ w) ◊ (T ◊ (T ◊ w)))) := by
416      nth_rw 1 [←lemma7]
417      - = ((x ◊ ((y ◊ x) ◊ x)) ◊ z) := by
418      nth_rw 1 [←lemma7]

419
420 have lemma9 (x y z : G) :
421     (x ◊ ((y ◊ ((z ◊ y) ◊ y)) ◊ x)) = x := by
422     calc
423         (x ◊ ((y ◊ ((z ◊ y) ◊ y)) ◊ x)) =
424             (x ◊ (((y ◊ ((z ◊ y) ◊ y)) ◊ x) ◊ ((y ◊ ((z ◊ y) ◊ y)) ◊ x))) := by
425             nth_rw 1 [lemma8]
426             - = (x ◊ (((y ◊ ((z ◊ y) ◊ y)) ◊ x) ◊ (((y ◊ ((z ◊ y) ◊ y)) ◊ x) ◊
427                 ((y ◊ ((z ◊ y) ◊ y)) ◊ x))) := by
428                 nth_rw 2 [lemma8]
429                 - = x := by
430                 nth_rw 1 [←op_law]

431
432 show _ by
433     intros x y z
434     calc
435         x = x ◊ ((x ◊ ((y ◊ x) ◊ x)) ◊ x) := by
436             nth_rw 1 [lemma9]
437             - = (x ◊ ((x ◊ ((y ◊ x) ◊ x)) ◊ x)) ◊ ((y ◊ (z ◊ (x ◊ z))) ◊
438                 ((x ◊ ((y ◊ x) ◊ x)) ◊ (y ◊ (z ◊ (x ◊ z))))) := by
439                 nth_rw 1 [←lemma5]
440                 - = x ◊ ((y ◊ (z ◊ (x ◊ z))) ◊ ((x ◊ ((y ◊ x) ◊ x)) ◊
441                     (y ◊ (z ◊ (x ◊ z))))) := by
442                     nth_rw 1 [lemma9]
443                     - = x ◊ (y ◊ (z ◊ (x ◊ z))) := by
444                     nth_rw 1 [lemma9]
```

7 Experiments on Other Equational Proofs

To assess the general potential of our approach, we evaluated our tool on a set of equational theorems obtained from the Equational Theories Project repository [6]. We selected all problems in the 13 Lean files `Proofs1` to `Proofs13` that have a proof and translated them to TPTP problem files, yielding 1431 benchmarks for our evaluation. One of them is Tao's challenge theorem 650 \Rightarrow 448.

For each file, we invoked our tool's TPTP problem generator, which parses the Lean theorems and produces corresponding TPTP problem files. For each problem, our tool was given 2700 seconds to produce a minimized proof using Vampire to find the initial proof and Vampire and Twee to find subproofs; on failure, the initial Vampire proof was output. A time limit of 10 seconds was used for each prover invocation. The experiments were conducted on a server equipped with a dual-socket AMD EPYC 9965 system (384 cores, 768 threads) running at 2.25–3.70 GHz with 3 TiB of DDR5 ECC RAM, and running Debian GNU/Linux 13 (kernel 6.17.13+deb13-amd64).

Overall, proofs for the 13 Lean files have an average length of 6.6 steps before minimization and 4.5 steps after minimization using the combination of small-step and abstracted problems and both provers. This corresponds to a 31.5% decrease, showing that even short proofs can often be made shorter.

Since longer proofs present more opportunities for minimization, we now focus on problems whose initial proofs have at least 15 steps. Table 1 compares proof lengths before and after minimization. The “Avg. before” column shows the average number of inference steps in the initial Vampire proofs. The remaining columns report the average proof length after minimization under four configurations, which differ in which problem variants are used to generate candidate lemmas: “BA” denotes the combination of the big-step and abstracted variants; “SA” denotes the combination of the small-step and abstracted variants; “BS” denotes the combination of the big- and small-step variants; and “BSA” denotes the combination of all three variants.

The results show an often substantial reduction in proof length. SA generally yielded the shortest proofs. Across all problems for the 13 Lean files, the average

Table 1. Comparison of proof lengths before and after minimization for problems with initial proofs of at least 15 steps

File	Num. problems	Avg. before	Avg. after			
			BA	SA	BS	BSA
Proofs1	135	17.3	16.0	13.1	13.3	13.3
Proofs2	117	16.9	14.4	11.5	11.5	11.5
Proofs3	108	19.3	15.6	10.9	10.7	10.9
Proofs4	125	19.1	14.3	10.9	11.4	11.4
Proofs5	116	20.1	17.6	12.8	11.9	11.9
Proofs6	115	25.6	18.9	12.5	12.6	12.6
Proofs7	117	37.2	19.7	11.8	11.9	11.9
Proofs8	114	24.4	15.6	12.3	13.1	13.1
Proofs9	112	39.8	29.0	13.1	14.1	14.1
Proofs10	101	21.5	16.0	8.0	11.0	11.0
Proofs11	110	25.4	22.4	13.0	14.0	14.0
Proofs12	123	24.6	16.5	8.0	8.5	8.5
Proofs13	38	35.3	27.7	9.1	10.1	10.1

476 reduction with SA is 56.7%. BS and BSA also produced substantial reductions,
 477 whereas BA generally yielded the least improvements.

478 It might seem counterintuitive that SA, which does not consider big-step
 479 problems in its search for the shortest proof of the main theorem, outperforms
 480 BSA. However, the nonmonotonicity is to be expected. Provers are nondeterministic,
 481 especially when invoked with a time limit. More importantly, our approach
 482 makes different heuristic choices when constructing the three proof segments de-
 483 pending on which problem variants are used. As a result, SA might find a short
 484 proof that escapes BSA.

485 The reduction in proof length is especially noticeable in individual cases. The
 486 problem $2666 \implies 3460$ has a Vampire proof with 51 inference steps, which our
 487 tool reduces to only 12 single rewrite steps, and $2923 \implies 2628$ is reduced from
 488 180 steps to only 34. The problem $3569 \implies 3957$ is reduced from 92 to 23 steps
 489 and, even more dramatically, $3957 \implies 3971$ is reduced from 141 steps to only 23.
 490 Furthermore, $2860 \implies 2660$ is reduced from 44 to 14 steps, and $723 \implies 872$ goes
 491 from 57 to 13 steps. Finally, $947 \implies 3897$ underwent the largest reduction, from
 492 151 to 10 steps. Overall, these results demonstrate that our approach produces
 493 shorter proofs across a diverse set of equational theorems.

494 8 Related Work

495 At least two other researchers took on Tao’s challenge. Kinyon [15] found a 24-
 496 step proof (excluding preprocessing) of $650 \implies \forall x, y. x = x \diamond y$ using Prover9
 497 [20], from which $650 \implies 448$ follows by instantiation. Later, Le Floch [10] de-
 498 veloped a pen-and-paper proof and translated it to Lean. The Lean proof relies
 499 on only 14 rewrite steps but includes additional reasoning as proof terms, and
 500 two of the rewrite steps are parallel, so the overall length is similar to ours. The
 501 proof idea is “loosely based” on the output of multiple Prover9 runs “with in-
 502 termediate results thrown in as assumptions or as goals”—in essence, a manual
 503 approximation of our approach.

504 We are aware of little work on automated proof minimization. Stachniak [24]
 505 designed an algorithm for constructing resolution proofs in propositional logics
 506 known as strongly finite logics. Amjad [2] and Cotton [9] introduced techniques
 507 for minimizing propositional resolution proofs. Gu et al. [13] developed Proof-
 508 Optimizer, which uses large language models to simplify Lean proofs.

509 Various SAT (satisfiability) and SMT (satisfiability modulo theories) solvers
 510 can minimize the number of axioms needed for a proof, but such minimization
 511 can yield longer proofs. In SAT solving, it is common to interleave search, which
 512 can be expressed as resolution steps, with formula-rewriting techniques that
 513 go beyond resolution. This interleaving, known as inprocessing [14], is highly
 514 effective and often yields both faster solving times and shorter proofs than either
 515 approach in isolation.

516 The idea of automatically mixing and matching proofs is not new. Sutcliffe
 517 et al. [26] introduced a method for combining automatically generated proofs to
 518 generate new ones. Their proofs are represented as DAGs, enabling the identifi-

519 cation and replacement of subproofs across different proofs. Proof combination
 520 is guided by heuristics that measure structural similarity, and a greedy search
 521 strategy is used to explore alternative combinations that yield proofs differing
 522 from the originals. In contrast to our approach, the main objective is to increase
 523 proof diversity rather than minimize proof length.

524 9 Conclusion

525 Historically, more research has gone into finding proofs automatically than into
 526 improving and presenting them. We introduced an approach for minimizing equa-
 527 tional proofs by mixing and matching the output of separate runs of Vampire
 528 and Twee, and implemented it in a new tool, Krympa. We used the tool to min-
 529 imize the proof of problem 650 \implies 448 from the Equational Theories Project
 530 from 62 to 20 steps, thereby providing a fully automatic solution to a challenge
 531 posed by Tao. We also obtained remarkable reductions on other problems orig-
 532 inating from the project. The shorter proofs are arguably easier to understand
 533 by humans and sometimes more general. Our work shows that proof automation
 534 and readability can go hand in hand.

535 Our approach could be extended in several ways. First, it could be generalized
 536 to support full first- or higher-order logic. Second, alternative lemma abstraction
 537 strategies could be explored. Third, proofs with more than three segments could
 538 be synthesized. Fourth, we might want to consider not only the number of steps
 539 but also term size when measuring proofs, as suggested by Le Floch [11]. Fifth,
 540 we could try to translate Vampire's superposition steps to Twee's structured
 541 equality chain format.

542 Some possible extensions specifically concern the implementation. Since proof
 543 generation relies heavily on external provers, performance could benefit from
 544 better scheduling of prover invocations, using adaptive time limits. Moreover, as
 545 the number of possible lemma combinations grows rapidly, exploiting parallelism
 546 at multiple levels—such as lemma re-proving, dependency graph construction,
 547 and proof construction—would be a natural extension of the current architecture.

548 **Acknowledgments.** We thank Bruno Le Floch, Andrew Reynolds, Stephan
 549 Schulz, and Uwe Waldmann for fruitful discussions. We thank Mark Summerfield
 550 for helpful textual suggestions.

551 Blanchette's research was cofunded by the European Union (ERC, Nekoka,
 552 101083038). Views and opinions expressed are however those of the authors only
 553 and do not necessarily reflect those of the European Union or the European
 554 Research Council. Neither the European Union nor the granting authority can
 555 be held responsible for them.

556 Heule's research is supported by the NSF under grant DMS-2434625 and
 557 funding from AFRL and DARPA under Agreement FA8750-24-9-1000.

558

References

- 559 1. The Lean Language Reference (2025), <https://lean-lang.org/doc/reference/latest/>
- 560 2. Amjad, H.: Compressing propositional refutations. In: Merz, S., Nipkow, T. (eds.) AVoCS 2006. Electronic Notes in Theoretical Computer Science, vol. 185, pp. 3–15. Elsevier (2006)
- 561 3. Bachmair, L., Dershowitz, N., Plaisted, D.A.: Completion without failure. In: Aït-Kaci, H., Nivat, M. (eds.) Rewriting Techniques, pp. 1–30. Academic Press (1989)
- 562 4. Bachmair, L., Ganzinger, H.: Strict basic superposition. In: Kirchner, C., Kirchner, H. (eds.) CADE 1998. LNCS, vol. 1421, pp. 160–174. Springer (1998)
- 563 5. Blanchette, J.C., Böhme, S., Fleury, M., Smolka, S.J., Steckermeier, A.: Semi-intelligible Isar proofs from machine-generated proofs. J. Automated Reas. **56**, 155–200 (2016)
- 564 6. Bolan, M., Breitner, J., Brox, J., Carlini, N., Carneiro, M., van Doorn, F., Dvorak, M., Goens, A., Hill, A., Husum, H., Mejia, H.I., Kocsis, Z.A., Floch, B.L., Bar-on, A.L., Luccioli, L., McNeil, D., Meiburg, A., Monticone, P., Nielsen, P., Osazuwa, E.O., Paolini, G., Petracci, M., Reinke, B., Renshaw, D., Rossel, M., Roux, C., Scanvic, J., Srinivas, S., Tadipatri, A.R., Tao, T., Tsyrklevich, V., Vaquerizo-Villar, F., Weber, D., Zheng, F.: The Equational Theories Project (2025)
- 565 7. Buch, A., Hillenbrand, T.: WALDMEISTER: Development of a high performance completion-based theorem prover (1996)
- 566 8. Clune, J., Qian, Y., Bentkamp, A., Avigad, J.: Duper: A proof-producing superposition theorem prover for dependent type theory. In: Bertot, Y., Kutsia, T., Norrish, M. (eds.) ITP 2024. LIPIcs, vol. 309, pp. 1–20. Leibniz-Zentrum für Informatik (2024)
- 567 9. Cotton, S.: Two techniques for minimizing resolution proofs. In: Strichman, O., Szeider, S. (eds.) Theory and Applications of Satisfiability Testing - SAT 2010, 13th International Conference, SAT 2010, Edinburgh, UK, July 11–14, 2010. Proceedings. Lecture Notes in Computer Science, vol. 6175, pp. 306–312. Springer (2010)
- 568 10. Floch, B.L.: Zulip post in “Machine Learning for Theorem Proving: A (Semi)-Autoformalization Challenge ($650 \rightarrow 448$)”. Zulip (2025), available at <https://leanprover.zulipchat.com/#narrow/channel/219941-Machine-Learning-for-Theorem-Proving/topic/A.20.28semi.29-autoformalization.20challenge.3A.20650.3D.3E448/near/518970125>
- 569 11. Floch, B.L.: Zulip post in “Machine Learning for Theorem Proving: A (Semi)-Autoformalization Challenge ($650 \rightarrow 448$)”. Zulip (2025), available at <https://leanprover.zulipchat.com/#narrow/channel/219941-Machine-Learning-for-Theorem-Proving/topic/A.20.28semi.29-autoformalization.20challenge.3A.20650.3D.3E448/near/568313777>
- 570 12. Geoff Sutcliffe: Stepping stones in the TPTP World. In: Benzmüller, C., Heule, M., Schmidt, R. (eds.) IJCAR 2024. pp. 30–50. LNCS (2024)
- 571 13. Gu, A., Piotrowski, B., Gloeckle, F., Yang, K., Markosyan, A.H.: ProofOptimizer: Training language models to simplify proofs without human demonstrations. CoRR [abs/2510.15700](https://arxiv.org/abs/2510.15700) (2025)
- 572 14. Järvisalo, M., Heule, M., Biere, A.: Inprocessing rules. In: Gramlich, B., Miller, D., Sattler, U. (eds.) IJCAR 2012. LNCS, vol. 7364, pp. 355–370. Springer (2012)
- 573 15. Kinyon, M.: Zulip post in “Machine Learning for Theorem Proving: A (Semi)-Autoformalization Challenge ($650 \rightarrow 448$)”. Zulip (2025),

- available at <https://leanprover.zulipchat.com/#narrow/channel/219941-Machine-Learning-for-Theorem-Proving/topic/A.20.28semi.29-autoformalization.20challenge.3A.20650.3D.3E448/near/518961204>
- 607 16. Knuth, D.E., Bendix, P.B.: Simple word problems in universal algebras. In: Leech,
608 J. (ed.) Computational Problems in Abstract Algebra, pp. 263–297. Pergamon
609 Press (1970)
 - 610 17. Kondylidou, L., Blanchette, J., Heule, M.: Tao's equational proof challenge ac-
611 cepted. Zenodo (2026), <https://doi.org/10.5281/zenodo.18624123>
 - 612 18. Kovács, L., Voronkov, A.: First-order theorem proving and Vampire. In: Sharygina,
613 N., Veith, H. (eds.) CAV 2013. LNCS, vol. 8044, pp. 1–35. Springer (2013)
 - 614 19. Limpert, J., From, A.H.: Aesop: White-box best-first proof search for Lean. In:
615 CPP 2023. pp. 253–266. Association for Computing Machinery (2023)
 - 616 20. McCune, W.: Prover9 and Mace4 (2005–2010)
 - 617 21. de Moura, L., Ullrich, S.: The Lean 4 theorem prover and programming language.
618 In: Bertot, Y., Kutsia, T., Norrish, M. (eds.) CADE 2021. LNCS, vol. 12699, pp.
619 625–635. Springer (2021)
 - 620 22. Norman, C., Avigad, J.: Canonical for automated theorem proving in Lean. In:
621 ITP 2025. LIPIcs, vol. 352, pp. 1–20. Leibniz-Zentrum für Informatik (2025)
 - 622 23. Smallbone, N.: Twee: An equational theorem prover. In: Platzer, A., Sutcliffe, G.
623 (eds.) CADE 2021. LNCS, vol. 12699, pp. 602–613. Springer (2021)
 - 624 24. Stachniak, Z.: Minimization of resolution proof systems. Fundam. Informaticae
625 **14**(1), 129–146 (1991)
 - 626 25. Sutcliffe, G.: The 12th IJCAR Automated Theorem Proving System
627 Competition—CASC-J12. AI Communications **38**, 3–20 (2025)
 - 628 26. Sutcliffe, G., Chang, C., McGuinness, D., Lebo, T., Ding, L., da Silva, P.P.: Com-
629 bining proofs to form different proofs. In: Fontaine, P., Stump, A. (eds.) PxTP
630 2011. pp. 60–73. LNCS (2011)
 - 631 27. Tao, T.: Machine learning for theorem proving: A (semi)-autoformalization
632 challenge (650 → 448). [https://leanprover-community.github.io/archive/
633 stream/219941-Machine-Learning-for-Theorem-Proving/](https://leanprover-community.github.io/archive/stream/219941-Machine-Learning-for-Theorem-Proving/), Lean Zulip thread,
634 created May 16, 2025
 - 635 28. Zhu, T., Clune, J., Avigad, J., Jiang, A.Q., Welleck, S.: Premise selection for a
636 Lean hammer. arXiv preprint arXiv:2506.07477 (2025)
 - 637
 - 638
 - 639